

Dingyi Duan

Module 6 Assignment

ADS-500B

04/11/2021

Module 6 Assignment Questions

Note that the answers to each of these questions should be the direct result of running appropriate Python or R code and not involve any manual processing of dataset files. Answers without either the code or results will not receive any grade.

1. For the next exercise, you are going to use the “airline_costs.csv” dataset.

The dataset has the following attributes:

- i. Airline name
- ii. Length of flight in miles
- iii. Speed of plane in miles per hour
- iv. Daily flight time per plane in hours
- v. Customers served in 1000s
- vi. Total operating cost in cents per revenue ton-mile
- vii. Revenue in tons per aircraft mile
- viii. Ton-mile load factor
- ix. Available capacity
- x. Total assets in \$100,000s
- xi. Investments and special funds in \$100,000s
- xii. Adjusted assets in \$100,000s

(Implement this exercise in Python language; import “pandas”, “statsmodels.api” libraries) Use a linear regression model to predict the number of customers each airline serves from its length of flight and daily flight time per plane.

```
# Q6.1

import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt

# Set max display
pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)

# Read the tsv file
df = pd.read_csv('C:/Users/DDY/Desktop/2021-Spring-textbooks/ADS-500B/Module6/airline_costs.csv', header=0)

print (df)

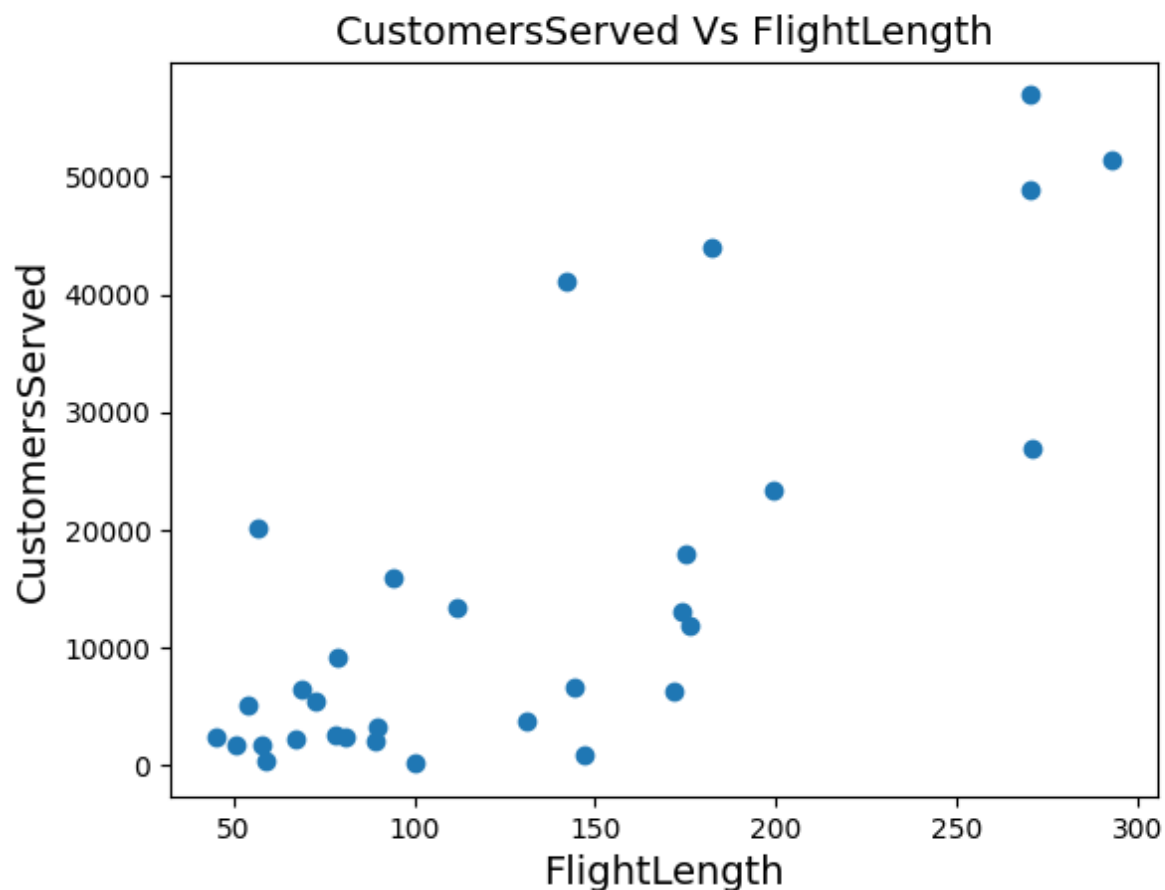
# Check for nulls for preprocessing
print (df.isna().sum())
```

```
Airline      0
FlightLength  0
PlaneSpeed   0
DailyFlightTime  0
CustomersServed  0
TotalOperatingCost  0
Revenue      0
LoadFactor   0
AvailableCapacity  0
TotalAssets  0
Investments  0
AdjustedAssets  0
dtype: int64
```

```
# Check for Multilinearity between 'DailyFlightTime' and 'FlightLength'
print (round(df['DailyFlightTime'].corr(df['FlightLength']), 2))
```

Check for multicollinearity between two independent variables: 0.48 (weak)

```
# Check for linearity between dependent and independent variables
# Linearity between 'CustomersServed' and 'FlightLength'
print (round(df['CustomersServed']. corr(df['FlightLength']),2))
# Correlation coefficient: 0.79 (strong)
|
plt.scatter(df['FlightLength'], df['CustomersServed'])
plt.title('CustomersServed Vs FlightLength', fontsize=14)
plt.xlabel('FlightLength', fontsize=14)
plt.ylabel('CustomersServed', fontsize=14)
plt.show()
```

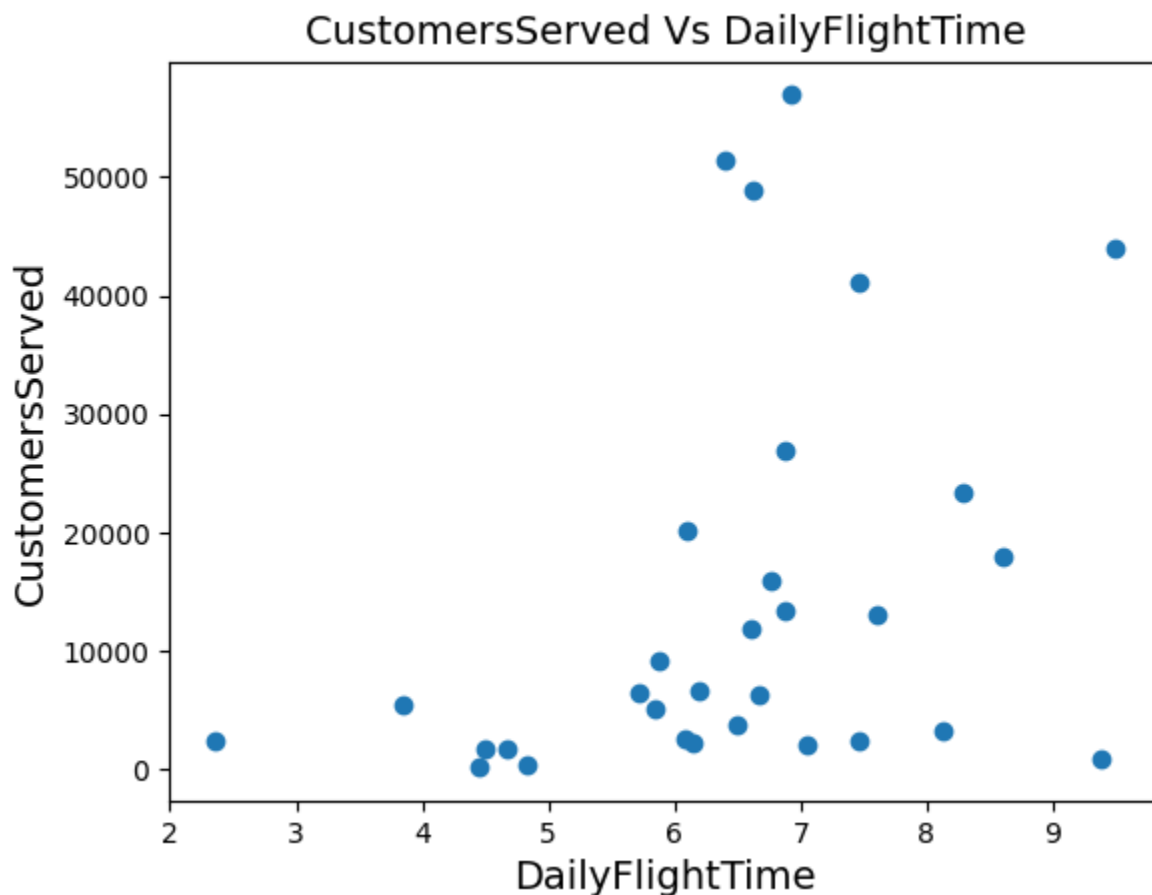


Correlation coefficients for 'CustomersServed' and 'FlightLength' and 'CustomersServed' and

'DailyFlightTime' : 0.79
 0.36

```
# Linearity between 'CustomersServed' and 'DailyFlightTime'
print (round(df['CustomersServed']. corr(df['DailyFlightTime']),2))
# Correlation coefficient: 0.36 (weak) as we see outliers

plt.scatter(df['DailyFlightTime'], df['CustomersServed'])
plt.title('CustomersServed Vs DailyFlightTime', fontsize=14)
plt.xlabel('DailyFlightTime', fontsize=14)
plt.ylabel('CustomersServed', fontsize=14)
plt.show()
```



```
# ===== Detect and clean up outliers =====

# Check outliers for 'FlightLength'
FlightLength_Q1 = df.FlightLength.quantile(0.25)
FlightLength_Q3 = df.FlightLength.quantile(0.75)
FlightLength_IQR = FlightLength_Q3 - FlightLength_Q1
FlightLength_out = (df.FlightLength < (FlightLength_Q1 - 1.5 * FlightLength_IQR)) | (df.FlightLength
> (FlightLength_Q3 + 1.5 * FlightLength_IQR))

print (FlightLength_out) # No outliers
```

```
0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9    False
10   False
11   False
12   False
13   False
14   False
15   False
16   False
17   False
18   False
19   False
20   False
21   False
22   False
23   False
24   False
25   False
26   False
27   False
28   False
29   False
30   False
Name: FlightLength, dtype: bool
```

```
DailyFlightTime_Q1 = df.DailyFlightTime.quantile(0.25)
DailyFlightTime_Q3 = df.DailyFlightTime.quantile(0.75)
DailyFlightTime_IQR = DailyFlightTime_Q3 - DailyFlightTime_Q1
DailyFlightTime_out = (df.DailyFlightTime < (DailyFlightTime_Q1 - 1.5 * DailyFlightTime_IQR)) | (df.DailyFlightTime
> (DailyFlightTime_Q3 + 1.5 * DailyFlightTime_IQR))

print(DailyFlightTime_out) # Outliers: row10, 28, 29

df = df[~(DailyFlightTime_out)] # Remove outliers from df
```

```

0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9    False
10   True
11   False
12   False
13   False
14   False
15   False
16   False
17   False
18   False
19   False
20   False
21   False
22   False
23   False
24   False
25   False
26   False
27   False
28   True
29   True
30   False
Name: DailyFlightTime, dtype: bool

```

Multiple Linear Regression Model:

```

# Perform Multiple Linear Regression using statsmodels
X = df[['FlightLength', 'DailyFlightTime']]
y = df['CustomersServed']

X = sm.add_constant(X) # adding a constant

model = sm.OLS(y, X).fit()
predictions = model.predict(X)

print_model = model.summary()
print(print_model)

# Equation: CustomersServed = 179.69 * FlightLength - 244.11 * DailyFlightTime - 7372.77
# R-squared = 0.654

```

Multiple Linear model:

$$\text{CustomersServed} = 179.69 * \text{FlightLength} - 244.11 * \text{DailyFlightTime} - 7372.77$$

$$\text{R-squared} = 0.654$$

OLS Regression Results						
Dep. Variable:	CustomersServed	R-squared:	0.654			
Model:	OLS	Adj. R-squared:	0.626			
Method:	Least Squares	F-statistic:	23.64			
Date:	Sun, 11 Apr 2021	Prob (F-statistic):	1.73e-06			
Time:	16:26:04	Log-Likelihood:	-296.16			
No. Observations:	28	AIC:	598.3			
Df Residuals:	25	BIC:	602.3			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-7372.7739	1.08e+04	-0.682	0.502	-2.96e+04	1.49e+04
FlightLength	179.6901	29.009	6.194	0.000	119.945	239.436
DailyFlightTime	-244.1079	1844.128	-0.132	0.896	-4042.160	3553.944
Omnibus:	4.032	Durbin-Watson:	1.876			
Prob(Omnibus):	0.133	Jarque-Bera (JB):	2.736			
Skew:	0.751	Prob(JB):	0.255			
Kurtosis:	3.301	Cond. No.	859.			

Next, build another regression model to predict the total assets of an airline from the customers served by the airline.

Simple Linear Regression Model:

```
# ===== Second linear regression model =====
# Linearity between 'CustomersServed' and 'FlightLength'
print(round(df['TotalAssets'].corr(df['CustomersServed']),2))
# Correlation coefficient: 0.89 (strong)

plt.scatter(df['CustomersServed'], df['TotalAssets'])
plt.title('TotalAssets Vs CustomersServed', fontsize=14)
plt.xlabel('TotalAssets', fontsize=14)
plt.ylabel('CustomersServed', fontsize=14)
plt.show()
```

Correlation coefficients for 'CustomersServed' and 'FlightLength' is 0.89

```

# Check outliers for 'CustomersServed'
CustomersServed_Q1 = df.CustomersServed.quantile(0.25)
CustomersServed_Q3 = df.CustomersServed.quantile(0.75)
CustomersServed_IQR = CustomersServed_Q3 - CustomersServed_Q1
CustomersServed_out = (df.CustomersServed
                        < (CustomersServed_Q1 - 1.5 * CustomersServed_IQR)) | (df.CustomersServed
                                                                              > (CustomersServed_Q3 + 1.5 * CustomersServed_IQR))

print (CustomersServed_out) # Outliers: row1,24,25


df = df[~(CustomersServed_out)] # Remove outliers from df

```

```

0.89
0    False
1     True
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9    False
11   False
12   False
13   False
14   False
15   False
16   False
17   False
18   False
19   False
20   False
21   False
22   False
23   False
24     True
25     True
26   False
27   False
30   False
Name: CustomersServed, dtype: bool

```

```
# Perform Multiple Linear Regression
x = df[['CustomersServed']]
y = df['TotalAssets']

x = sm.add_constant(x) # adding a constant

model = sm.OLS(y, x).fit()
predictions = model.predict(x)

print_model = model.summary()
print(print_model)

# Equation: CustomersServed = 0.0072 * FlightLength + 2.86
# R-squared = 0.472
```

Linear model:

$$\text{TotalAssets} = 0.0072 * \text{CustomersServed} + 2.86$$

$$\text{R-squared} = 0.472$$

OLS Regression Results						
Dep. Variable:	TotalAssets	R-squared:	0.472			
Model:	OLS	Adj. R-squared:	0.449			
Method:	Least Squares	F-statistic:	20.58			
Date:	Sun, 11 Apr 2021	Prob (F-statistic):	0.000148			
Time:	16:40:36	Log-Likelihood:	-143.19			
No. Observations:	25	AIC:	290.4			
Df Residuals:	23	BIC:	292.8			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	2.8593	21.936	0.130	0.897	-42.519	48.238
CustomersServed	0.0072	0.002	4.536	0.000	0.004	0.010
Omnibus:	22.944	Durbin-Watson:	2.353			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	40.383			
Skew:	1.764	Prob(JB):	1.70e-09			
Kurtosis:	8.131	Cond. No.	1.96e+04			
Notes:						
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.						
[2] The condition number is large, 1.96e+04. This might indicate that there are strong multicollinearity or other numerical problems.						

Do you have any insight about the data from the last two regression models? (20 points)

Answer: For multiple linear regression models, we are dealing with multiple variables that may have weak linear relationships with our predictor. From the first summary table we can see that the p value for 'DailyFlightTime' is 0.896 which means this variable is not statistically significant, which also matches the weak correlation coefficient 0.36. However, when we finish building the model, we obtain a better R-squared value of 0.654; Compared to the simple linear regression model where we have the dependent and independent variables are highly correlated to each other, we only obtain an R-squared value of 0.472 which leads to a less efficient model than multiple linear one. To conclude, it is more accurate to build a linear regression model with multiple variables than a single independent variable.

2. For this clustering exercise, you are going to use the data on women professional golfers performance on the LPGA, 2008 tour ('lpga2008.csv' dataset). The dataset has the following attributes:

- i. Golfer: name of the player
- ii. Average Drive distance
- iii. Fairway Percentage
- iv. Greens in regulation: in percentage
- v. Average putts per round
- vi. Sand attempts per round
- vii. Sand saves: in percentage
- viii. Total Winnings per round
- ix. Log: Calculated as (Total Win/Round)
- x. Total Rounds
- xi. Id: Unique ID representing each player

(Implement this exercise in R language; import 'cluster' library)

Use agglomerative clustering and divisive clustering on this dataset to find out which players have similar performance in the same season. Visualize the clusters using dendrograms for both types of clustering models. (20 points)

```
# Q6.2

library(cluster)
library(factoextra)
library(NbClust)

# Read the file
lpga = read.table('C:/Users/DDY/Desktop/2021-spring-textbooks/ADS-500B/Module6/lpga2008.csv', header=T, sep=',')

# Remove '?' and '0's
lpga[lpga == '?' && '0'] = NA
lpga = na.omit(lpga)

# Count nulls >>> no nulls
sum(is.na(lpga))
```

```
> # Count nulls >>> no nulls
> sum(is.na(lpga))
[1] 0
> |
```

```
# Standardizing the dataset
newlpga = lpga[-1]
row.names(newlpga)=lpga$Golfer
newlpga$Id = NULL
view(newlpga)
lpga_st = scale(newlpga)
```

```
# Determining 'k'
```

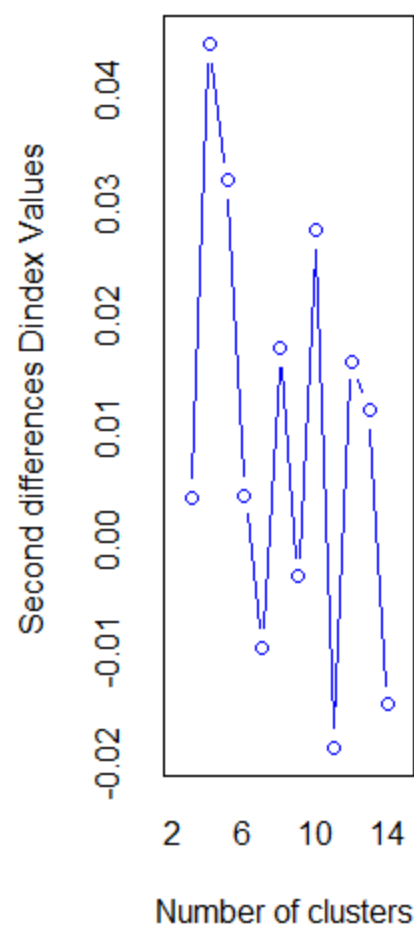
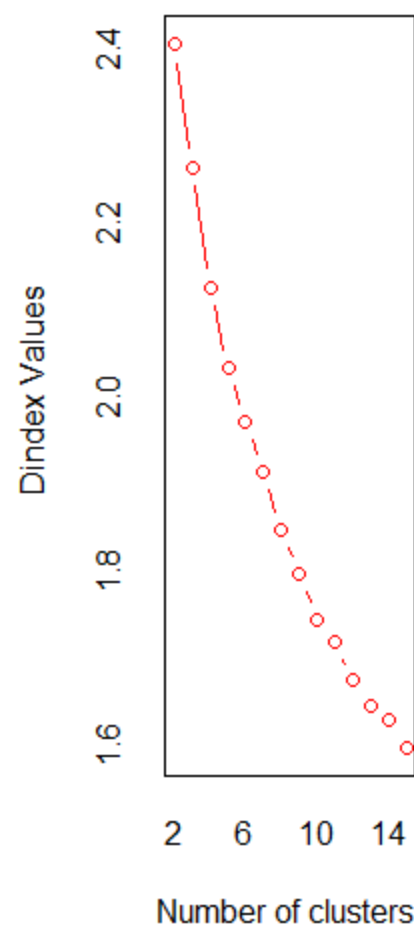
```
# using nbClust() to determine 'k' value
NbClust(lpga_st,distance='euclidean', method='kmeans') # According to majority rule,
# the best number of clusters is 2
```

```
*****
```

```
* Among all indices:
* 7 proposed 2 as the best number of clusters
* 3 proposed 3 as the best number of clusters
* 3 proposed 4 as the best number of clusters
* 4 proposed 5 as the best number of clusters
* 1 proposed 7 as the best number of clusters
* 1 proposed 10 as the best number of clusters
* 3 proposed 14 as the best number of clusters
* 1 proposed 15 as the best number of clusters
```


```
***** conclusion *****
```

```
* According to the majority rule, the best number of clusters is 2
```



```
# ===== Agglomerative (Agnes) clustering =====  
# Build the clusters in the Agglomerative method using the agnes() function  
aclusters = agnes(lpga_st,metric = 'euclidean',stand=FALSE)  
# Agglomerative coefficient  
aclusters$ac  
  
> # Agglomerative coefficient  
> aclusters$ac  
[1] 0.8088495  
< |
```

pg_a_st
agnes (? "average")



Dendrogram of Agnes

```
# Plotting Agnes Dendrogram
pltree(aclusters, cex=0.6, hang=-1, main = 'Dendrogram of Agnes')

# Indicating rectangles on the plot to visualize 2 clusters
rect.hclust(aclusters, k=2, border = 2:5)


# Grouping clusters using cutree() function
grp_agnes = cutree(aclusters, k=2)

# Forming table to see the cluster size
table(grp_agnes)

# visualization using fviz_cluster
fviz_cluster(list(data=lpga_st, cluster=grp_agnes))
```



```
> table(grp_agnes)
grp_agnes
  1    2
156   1
```




```
# ===== Divisive (Diana) clustering =====
# Build the clusters in the Divisive method using the diana() function
diclusters = diana(lpga_st,metric = 'euclidean',stand=FALSE)
# Divisive coefficient
diclusters$dc
```

```
> # Divisive coefficient
> diclusters$dc
[1] 0.8719543
> |
```

```
# Plotting Diana Dendrogram
pltree(diclusters, cex=0.6, hang=-1, main = 'Dendogram of Diana')

# Indicating rectangles on the plot to visualize 2 clusters
rect.hclust(diclusters, k=2, border = 3:5)
```



```
# Grouping clusters using cutree() function
grp_diana = cutree(diclusters, k=2)

# Forming table to see the cluster size
table(grp_diana)

# visualization using fviz_cluster
fviz_cluster(list(data=lpga_st, cluster=grp_diana))
```

```
> table(grp_diana)
grp_diana
 1  2
87 70
```

